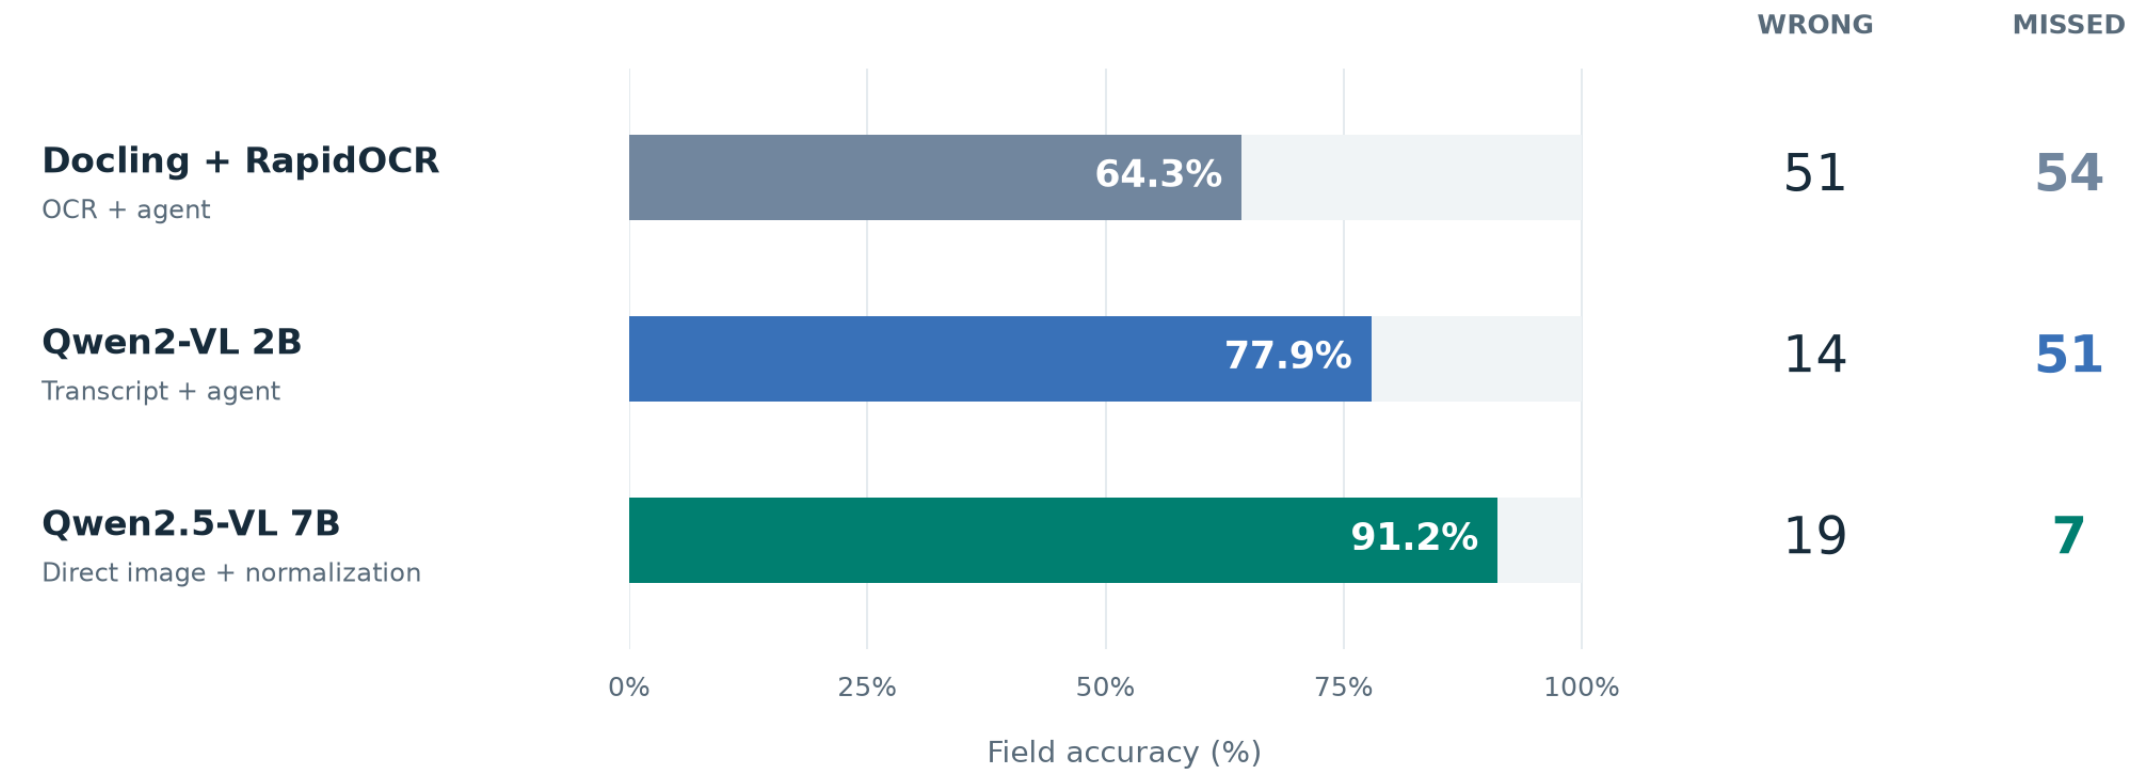
Better extraction, fewer missing fields

Same 14 gold pages / 294 populated values / Current evaluation rubric

+26.9 percentage points

Docling pipeline to direct Qwen2.5-VL extraction



Accuracy = correct / (correct + wrong + missed).

14 extra year predictions in the final run are flagged separately; they are outside this accuracy denominator.

Source: saved gold_eval_summary.json files, July 5-7, 2026. Small development set; not a held-out test.

Where the gains came from

Six selected fields / 14 populated values per field / Accuracy (%)

	Docling + RapidOCR OCR + agent	Qwen2-VL 2B Transcript + agent	Qwen2.5-VL 7B Direct image + normalization
Marriage day	35.7%	57.1%	**92.9%**
Marriage month	14.3%	35.7%	**92.9%**
License number	64.3%	7.1%	**85.7%**
Officiant signature	35.7%	42.9%	**85.7%**
Marriage location	7.1%	28.6%	**42.9%**
Relationship	57.1%	64.3%	**42.9%**

Dates and signatures improve; location and relationship remain difficult.

Selection includes gains, a non-monotonic identifier result, and remaining weaknesses. Darker cells = higher accuracy.

Source: saved gold_eval_per_field.csv files, July 5-7, 2026. See docs/DATA_AND_EVALUATION.md for the full results.